

Reachability Search over Convex Transition Sets for Multi-Regrasp Manipulation Planning

Julia López Gómez and Steve Tonneau

School of Informatics, University of Edinburgh, United Kingdom

J.Lopez-Gomez@sms.ed.ac.uk, stonneau@ed.ac.uk

Abstract—A robot opening a bottle or turning a valve runs out of wrist range long before the object reaches its goal, and must grasp, turn, release and grasp again. How many times depends on the goal, the joint limits and the obstacles, and is not known before planning. Sampling-based manipulation planners cannot establish, after a finite computation, that a plan uses the fewest grasps or that none exists, and optimisation-based ones must fix the number of grasps in advance. We build the manipulation graph out of convex sets: intersecting a convex decomposition of the collision-free space of robot and object with a convex set of grasping configurations gives *transition sets*, the only places where the robot can grasp, turn and release. An A* searches them with nodes that pair a set with the interval of object configurations reachable inside it; reachability then propagates by interval arithmetic, and pruning is exact. The graph is built with linear programs alone, and the search returns a plan with the fewest grasps these sets allow, or reports that none exists, without fixing that number in advance. On a 10-DOF cap-turning task with a Franka Panda, where our earlier mixed-integer formulation is unreliable, it plans in 4ms. **[TODO: headline comparison (E1–E4) and validated result (E5).]**

Index Terms—Manipulation planning, regrasping, reachability, graph search, convex decompositions.

I. INTRODUCTION

A robot opening a bottle, turning a valve or winding a crank cannot finish in a single grasp: the wrist runs out of range long before the object reaches its goal. It has to grasp, turn as far as it can, let go, rotate the wrist back and grasp again. Planning this means deciding where to grasp and release, how far to turn each time, how many times to regrasp, and how to move in between. None of these can be decided on its own. An obstacle beside the object shortens the usable part of the wrist range, which changes how far one grasp can turn the object, and with it the number of grasps the task needs.

What makes this hard is how thin the feasible sets are, and the fact that the task accumulates. For an arm and a hand, the configurations that hold an object at a fixed grasp satisfy several equality constraints on the pose of the hand, and so form a thin, curved set in a configuration space of ten dimensions or more. Only its collision-free part is usable, and that part lies precisely where the hand is closest to the obstacles. The object, meanwhile, moves only while held and keeps its angle once released, so what the robot can still reach depends on all it has done so far.

This work was supported by the UKRI AI Centre for Doctoral Training in Dependable and Deployable AI for Robotics (CDT-D2AIR) **[TODO: grant number].**

For a robot moving alone, convex decompositions of the collision-free space [1]–[4] made this kind of reasoning tractable: Graphs of Convex Sets return trajectories that are globally optimal relative to the decomposition [5]. None of that machinery decides when to grasp or release. Planners that do reason about regrasping either sample the manipulation graph [6]–[8], enumerate a finite set of grasps [9], [10], or fix the sequence of grasps and optimise the motions [11], [12]. Mixed-integer formulations over convex sets, our own included [13], are complete and optimal relative to the decomposition, but only for a number of grasps fixed beforehand, and the indicator constraints that select among the sets are poorly conditioned when those sets are thin (Section IV-G). After a finite computation, none of them can say that a plan uses the fewest grasps, or that no plan exists.

Footstep planning met the same obstacle and got past it by searching over sets instead of points: CASSR [14] runs A* on convex sets of footholds reachable from the start, planning 30 steps in under 125 ms. Regrasping asks for more than a change of domain: footholds come ready-made from the environment, whereas grasping configurations must be built from nonlinear kinematics, and a grasp, unlike a footstep, moves the object, so reachability has to carry the object’s state.

We build the manipulation graph of [6] out of convex sets. A convex decomposition of free space gives the *free sets*; a convex inner approximation of the grasping configurations, grown from a seed grasp along the null space of the grasp constraints [15], gives the *grasp set*; their intersections are the *transition sets*, the only places where the robot can grasp, turn and release the object. We search them with an A* whose nodes pair a set with the interval of object configurations reachable inside it. Convexity makes every point of a set reachable from any other along a line that never leaves it, and the interval makes propagation and pruning exact, so the search needs no sampling, no integer variables and no horizon fixed in advance. Our contributions are:

- *transition sets*, a convex representation of the manipulation graph, built for a 7-DOF arm and validated against a collision checker where the decomposition is least reliable (Sections III, IV-A–IV-C);
- a *reachability A** over them, complete and returning the fewest grasps relative to the sets, with a solver-free stage that recovers the motions (Sections IV-D–IV-F);
- an evaluation on cap-turning tasks, up to the 10-dimensional space of a Franka arm, hand and cap, against

our mixed-integer planner and a sampling-based one (Sections V–VI). **[TODO: update with E1–E5.]**

[CHECK: anonymise the self-citation if double-blind.]

II. RELATED WORK

A. Regrasp and manipulation planning

The first regrasp planners searched over finite sets of grasps and placements [16], [17]. The idea persists in regrasp graphs built over thousands of precomputed grasps [9], and in a bi-manual planner certified complete with respect to its discretisation [10]. Such planners are complete only with respect to the grasps they enumerate: a solution that needs a grasp between two discretised ones is missed. The manipulation graph [6] made grasps and placements continuous. It was extended to multi-modal planning [7], to asymptotically optimal sampling [18], and to constraint graphs explored with constrained samplers [8], [19]. Their guarantees are asymptotic: they hold in the limit of infinitely many samples. Task and motion planners add symbolic decisions on top [20]. In optimisation-based TAMP, an outer search enumerates skeletons with a fixed number of mode switches, and a local solver handles each one [11]. Closest in spirit, Liu et al. [21] plan contact-rich planar manipulation on a graph of reachable sets of object poses **[CHECK: characterisation]**. Our transition sets are the nodes of the manipulation graph of [6], built as convex sets rather than sampled. This is what lets us guarantee completeness and the minimum number of grasps after a finite computation, relative to the sets.

B. Contact planning over convex sets

Footstep planning faces the same combinatorics as regrasp-ing: a contact sequence whose length is not known in advance. A* over discretised footholds [22] is complete only at the resolution of the discretisation. Mixed-integer programs over convex terrain [23] avoid discretisation and are globally optimal, but the number of steps, or a bound on it, must be fixed, and their solve times are hard to bound. ℓ_1 relaxations [24] make them fast at the cost of these guarantees. **[CHECK: SLIM characterisation.]** CASSR [14] instead searches with A* over convex sets of footholds reachable from the start, expanded by Minkowski sums, and places the footsteps with a quadratic program once the sets are chosen. CASSR discretises the yaw of the feet, and inflates its heuristic for speed, so its optimality is empirical rather than guaranteed. We follow the same principle, but a node’s reachable set carries the configuration of the object rather than a foothold. In the coupled tasks we consider it is an interval, which lets us keep an admissible heuristic and exact pruning, with nothing discretised.

Graphs of Convex Sets (GCS) [5], [25] compute shortest paths over convex decompositions through a convex relaxation, and are globally optimal relative to the decomposition when the relaxation is tight. They plan the motion of a robot in one mode, with no switches of grasp and no object state carried from one switch to the next. They have been extended to contact-rich planar pushing [26]. Heuristic searches handle

graphs too large to solve at once [27]. GCS* is complete and optimal when it can decide exactly whether one partial path dominates another, i.e., whether its reachable set contains the other’s at lower cost. In general this check must be approximated [27]. With interval reachable sets, it is exact.

C. Convex decompositions of free space

Convex decompositions inner-approximate free space with convex sets. IRIS grows them by alternating separating hyperplanes and inscribed ellipsoids [1]. C-IRIS certifies regions as collision-free with sums-of-squares programs [2]. IRIS-NP, IRIS-NP2 and IRIS-ZO scale to more dimensions, with a probabilistic guarantee: the fraction of a region’s volume in collision is below a threshold, with a given confidence [3], [28]. Clique covers automate the seeding [4], and GPU implementations approach online rates [29]. A volume guarantee says nothing about sets of near-zero volume, such as grasp sets. This is why we validate the transition sets separately (Section IV-C).

III. PROBLEM FORMULATION

A. Transit, transfer and transition sets

A robot with configuration $x \in \mathcal{X} \subset \mathbb{R}^{n_r}$ (arm and fingers) manipulates an object with configuration α in a known static scene. The state $q = (x, \alpha)$ lives in the composite configuration space $\mathcal{C}$ [6], and $\mathcal{C}_{\text{free}} \subseteq \mathcal{C}$ is the set of states without collisions other than the contacts a grasp requires. The object cannot move on its own. In the *placement space* $\mathcal{C}_P$ it rests stably, and in the *grasp space* $\mathcal{C}_G$ the robot holds it. A manipulation path alternates *transit* segments in $\mathcal{C}_P$, where the robot moves alone and α is constant, and *transfer* segments in $\mathcal{C}_G$, where the robot moves the object at a fixed relative pose. The two can only meet where the object is at once grasped and stably placed, in $\mathcal{C}_G \cap \mathcal{C}_P \cap \mathcal{C}_{\text{free}}$. We call these the *transition sets*; they are the nodes of the manipulation graph [6].

B. Multi-regrasp tasks with a coupled object coordinate

We study the tasks captured by the following assumption, which include caps, knobs, valves and cranks.

Assumption 1: (i) The object has one coordinate $\alpha \in [\alpha_{\min}, \alpha_{\max}]$ and is always stably placed, so $\mathcal{C}_P = \mathcal{C}$. (ii) Holding the object couples α to one robot joint w , the *wrist*: during a transfer, $\dot{\alpha} = -\dot{w}$. The sign follows from the orientation of the axes. (iii) Neither the grasp constraints nor collisions depend on α : $\mathcal{C}_G = G \times [\alpha_{\min}, \alpha_{\max}]$ and $\mathcal{C}_{\text{free}} = \mathcal{X}_{\text{free}} \times [\alpha_{\min}, \alpha_{\max}]$ for some $G, \mathcal{X}_{\text{free}} \subseteq \mathcal{X}$.

Point (iii) holds for rotationally symmetric objects, such as the cap in our experiments. Point (i) has a consequence we rely on throughout: every grasping configuration is also a placement, so a transition set is simply a collision-free piece of the grasp space, $G \cap \mathcal{X}_{\text{free}}$ up to the object range. A transition set is therefore not only where the robot switches between the two kinds of motion, but also where it carries the object, so one set can host a grasp, a whole turn and the release that follows.

The wrist has a limited range $[w_{\min}, w_{\max}]$, so turning the object by more than this range requires several grasps. Without

loss of generality, let $\alpha_{\text{goal}} \leq \alpha_{\text{init}}$. Since α affects neither collisions nor grasps, a plan never gains from turning the object backwards, and the task works like a ratchet. A *stroke* grasps the object and drives the wrist forward, turning the object by the wrist travel; one stroke is one grasp, and counting strokes counts grasps. The robot then either lets go and rotates the wrist back, with α fixed, to begin another stroke in the same set — a *return* — or lets go and leaves the set — a *release*. Everything else is transit in $\mathcal{X}_{\text{free}}$ at fixed α .

C. Problem statement

Given $q_{\text{init}} = (x_{\text{init}}, \alpha_{\text{init}})$ and $q_{\text{goal}} = (x_{\text{goal}}, \alpha_{\text{goal}})$, find a continuous path from q_{init} to q_{goal} in $\mathcal{C}_{\text{free}}$, made of transit segments, strokes and returns, that uses the fewest grasps. Ties are broken by the number of changes of convex set along the path; path length is not optimised.

In terms of reachability, let $\mathcal{R}_k \subseteq \mathcal{C}_{\text{free}}$ be the set of states reachable from q_{init} with at most k grasps. The minimum number of grasps is the smallest k such that $q_{\text{goal}} \in \mathcal{R}_k$, and the goal is unreachable if there is no such k . Neither $\mathcal{R}_k$ nor $\mathcal{C}_{\text{free}}$ can be computed exactly. Our planner computes $\mathcal{R}_k$ exactly for an inner approximation of $\mathcal{X}_{\text{free}}$ and G by convex polytopes, and its guarantees hold relative to that approximation.

IV. METHOD

The planner has an offline stage, which depends only on the robot, the object and the scene, and an online stage that is run for each query. Offline, it builds the free sets (Section IV-A) and the grasp set (Section IV-B), and intersects them to obtain validated transition sets and the graph over them, using linear programs alone (Section IV-C). Online, the reachability A^* chooses the sequence of sets and grasps (Section IV-D), and a second stage turns it into motions (Section IV-E). New goals and start configurations reuse the offline stage unchanged. Fig. 1 illustrates the online stage on the task of Section V, which we use throughout as a running example.

A. Free sets

We inner-approximate $\mathcal{C}_{\text{free}}$ by polytopes $P_j = \{q : A_j q \leq b_j\}$, $j = 1, \dots, J$, computed with clique covers of visibility graphs [4] and IRIS-NP2 regions [28] in Drake [30]. By Assumption 1(iii), the object coordinate does not affect collisions. We therefore replace each region by its slice at $\alpha = 0$, $F_j = \{x : (x, 0) \in P_j\}$. This is sound: if $(x, 0)$ is collision-free, so is (x, α) for every α . The robot can then occupy any $x \in F_j$ at any object angle, and the planner tracks α separately.

B. Grasp set

A grasp fixes the pose of the hand frame E relative to the object frame C , a nonlinear function of the arm joints. We write the resulting conditions as

$$h(x) = 0, \quad g(x) \leq 0, \quad (1)$$

where h collects the equalities a valid grasp must satisfy, typically the components of ${}^C p^E(x)$ and of the hand axes that

the grasp pins down, and g the inequalities, such as tolerance bands on the remaining components, an approach direction and a finger opening. Section V instantiates them for our tasks. The set $\{h = 0, g \leq 0\}$ is thin and curved, so it cannot be written as a polytope. We approximate it locally, using two ideas from NLP sampling [15]. First, a seed grasp x_0 is found by Gauss–Newton minimisation of the constraint violation $\|[h(x); \max(g(x), 0)]\|^2$ from a random start. Second, the directions that preserve the grasp to first order are the null space of the Jacobian J_h of the equality constraints, which NLP sampling uses to move along constraint manifolds. The grasp set keeps the configurations that stay within ϵ of this null space, around x_0 :

$$G = \{x : |h_0 + J_h(x - x_0)| \leq \epsilon, \quad g_0 + J_g(x - x_0) \leq \epsilon, \\ x_{\min}^G \leq x \leq x_{\max}^G, \quad |x_j - x_{0,j}| \leq \rho \quad \forall j \in \mathcal{A}\}, \quad (2)$$


where h_0, J_h, g_0, J_g are evaluated at x_0 , ϵ is a tolerance, and $\mathcal{A}$ are the joints other than the wrist.

Two properties shape this set. First, whenever a joint's motion leaves h and g unchanged, its columns in J_h and J_g vanish and G spans that joint's entire range. This is what allows a stroke to run inside a single set, and it is exact rather than approximate: it holds for the wrist of any arm whose last axis passes through the origin of E along the approach direction, since rotating about it moves neither the position nor the approach axis of the hand. Second, J_h leaves a null space of dimension $n_r - \text{rank}(J_h)$ that the constraints of (2) bound only through the joint limits. Far from x_0 the linearisation is meaningless, so the trust region ρ keeps the set where the linearisation is valid; Section VI quantifies what happens without it. Seeds on other parts of the grasp manifold give further grasp sets that the search treats in the same way. We use a single seed in this paper (see Section VII).


C. Transition sets and the graph

For every free set with $F_k \cap G \neq \emptyset$, we form the transition set $T_k = G \cap F_k$. Two LPs give its extreme wrist values, $w_k^+ = \max\{w : x \in T_k\}$ and $w_k^- = \min\{w : x \in T_k\}$, together with points $x_k^+, x_k^- \in T_k$ attaining them. A stroke inside T_k turns the object by at most $L_k = w_k^+ - w_k^-$, and we discard sets with $L_k < L_{\min}$. Since G is thin in every direction but the wrist, T_k is a sliver along the wrist axis, and by convexity the whole segment $[x_k^-, x_k^+]$ lies inside it: the robot can sweep the entire wrist range of T_k without leaving it.

a) *Validation*: The free sets are collision-free only in a volume sense: IRIS-NP2 bounds the fraction of each region's volume in collision [28]. A transition set has almost no volume, so it can lie entirely in collision without violating that bound. We therefore validate each T_k on its own. Because it is a sliver, this is cheap: we test densely spaced configurations along $[x_k^-, x_k^+]$ with a collision checker, and keep only the wrist sub-intervals on which every test passes. Each of these is T_k intersected with a bound on w , and is therefore still a polytope. **[TODO: implement (E5); state the check resolution;**



(a) What the search tracks



(b) The same plan in the wrist-cap plane

Fig. 1. The reachability A^* on the running example, task T2 of Section V, query Q0: turn the cap from $\alpha = 3$ to -1 rad and bring the arm back to its start. (a) Each column is one node on the returned path: the set the robot is in (strip) and the interval of cap angles reachable there (bar). The black line is the cap angle along the final plan. Entering the transition set T_6 changes nothing; each stroke widens the interval by up to the wrist range of T_6 (2 rad); returns and transit keep it. Leaving T_6 towards F_0 needs some further wrist motion, which lowers the top of the last bar. The search stops at the first node whose set contains the goal arm configuration and whose interval contains α_{goal} . (b) While grasping, the cap turns as the wrist moves forward; returns move the wrist back with the cap fixed. Two strokes are the minimum, since each turns the cap by at most 2 rad. **[TODO: regenerate once the transition sets are validated (E5).]**

this is a dense check, not a certificate. Certifying these thin sets with C-IRIS [2] is an alternative.]

b) *Graph*: The graph has two kinds of edges. A *transit* edge joins F_i and F_j if they intersect, and its door point is the Chebyshev centre of $F_i \cap F_j$. A *switch* edge joins F_i and T_k if they intersect. For it, we store the wrist range $[a_{ik}^-, a_{ik}^+]$ of $F_i \cap T_k$ and the points attaining its ends. Each of these quantities is an LP feasibility check, a Chebyshev centre or an extreme point. No integer variables or vertex enumeration are needed, and the number of LPs is $O(J^2 + JK)$ for J free sets and K transition sets.

D. Reachability A^*

We use A^* in the same form and with the same syntax as [14], [22] (Algorithm 1). Only the node structure, the expansion, the check for nodes already expanded, the goal test and the costs are specific to our problem. They are highlighted in blue and described below.

1) *The A^* algorithm*: A^* keeps a priority queue ordered by the cost of the path to each node plus an estimate of the cost to the goal (line 15). It starts with one node per free set containing x_{init} , with the interval $\{\alpha_{\text{init}}\}$ (line 1). At each iteration, it takes the cheapest node (line 3) and skips it if an equivalent or better node was already expanded (line 4). If the node is a goal (line 6), the plan is read by following parents back to the start (line 17). Otherwise, the node is expanded (line 8), and its children are costed and queued (lines 9–16). If the queue empties first, no plan exists within the sets.

2) *Node structure*: As in CASSR, a node does not represent a configuration but a set reachable from the start (Struct 1). A node in a free set F_i with interval $[\alpha^-, \alpha^+]$ states that every $x \in F_i$ is reachable together with every $\alpha \in [\alpha^-, \alpha^+]$; since F_i is convex and α is constant during transit, a straight line joins any two of its points without leaving it. A node in a transition set T_k also records the wrist range $[a^-, a^+]$ through which T_k was entered, and a phase: *entered* (not yet grasping), *holding* (after a stroke) or *released* (after a return, with the wrist at w_k^-). The interval carries all the information about the path so

Algorithm 1 A^* search; the parts specific to our problem are in blue

```

1: addStartNodesToQueue() ▷ Sec. IV-D1
2: while hasNodesToCheck() do
3:   nodeToExpand = getCheapestNode()
4:   if nodeAlreadyExpanded() then ▷ Sec. IV-D4
5:     skipToNextNodeInQueue()
6:   if hasReachedTheGoal() or hasTimedOut() then
7:     stopSearch()
8:   childNodesToCheck = expandNode() ▷ Alg. 2
9:   for childNode ∈ childNodesToCheck do
10:    if nodeAlreadyExpanded(childNode) then
11:      continue ▷ skip node
12:    costOfEdge = getCostOfEdge() ▷ Sec. IV-D6
13:    addEdgeToGraph(childNode, costOfEdge)
14:    costToGoal = estimateCostToGoal() ▷ Eq. (3)
15:    nodeCost = costOfPath(childNode) + costToGoal
16:    addNodeToQueue(childNode, nodeCost)
17: plan = getBestPathToEndNode() ▷ NONE if queue empty

```

Struct 1 Node

```

parent: NODE* ▷ parent node
setId: INT ▷ free set  $F_i$  or transition set  $T_k$ 
interval:  $[\alpha^-, \alpha^+]$  ▷ reachable object configurations
phase: ENUM ▷ transition sets: entered, holding, released
entryRange:  $[a^-, a^+]$  ▷ transition sets: wrist range at entry

```

far: reaching F_3 before any grasp and after a stroke of 2 rad gives the intervals $\{3\}$ and $[1, 3]$, two different nodes.

3) *expandNode()*: Algorithm 2 lists the children of a node. A stroke turns the object by any amount up to c , the wrist travel left in T_k ($w_k^+ - a^-$ after entering, L_k after a return), so it lowers α^- and keeps α^+ . Since the object can only be grasped inside transition sets, strokes are the only moves that widen the interval. Leaving T_k towards F_j while holding the object is a release, and it can force some turning: the wrist

Algorithm 2 `expandNode`, for a node in set S with interval $[\alpha^-, \alpha^+]$, phase ϕ and entry range $[a^-, a^+]$

```

1: nodeList = []
2: if  $S = F_i$  is a free set then
3:   for all  $F_j$  with  $F_i \cap F_j \neq \emptyset$  do ▷ transit
4:     nodeList.add( $F_j$ ,  $[\alpha^-, \alpha^+]$ )
5:   for all  $T_k$  with  $F_i \cap T_k \neq \emptyset$  do ▷ enter
6:     nodeList.add( $T_k$ ,  $[\alpha^-, \alpha^+]$ , entered,  $[a_{ik}^-, a_{ik}^+]$ )
7: else ▷  $S = T_k$  is a transition set
8:   if  $\phi \in \{\text{entered}, \text{released}\}$  then ▷ stroke
9:      $c \leftarrow w_k^+ - a^-$  if  $\phi = \text{entered}$ , else  $c \leftarrow L_k$ 
10:     $\beta \leftarrow \max(\alpha^- - c, \alpha_{\text{goal}}, \alpha_{\text{min}})$ 
11:    nodeList.add( $T_k$ ,  $[\beta, \alpha^+]$ , holding)
12:   if  $\phi = \text{holding}$  then ▷ return
13:     nodeList.add( $T_k$ ,  $[\alpha^-, \alpha^+]$ , released)
14:   if  $\phi \in \{\text{entered}, \text{holding}\}$  then ▷ release or leave
15:     for all  $F_j$  with  $F_j \cap T_k \neq \emptyset$  do
16:        $d \leftarrow \max(0, a_{jk}^- - a^+)$  ▷ forced turning
17:       if  $\phi = \text{holding}$  or  $d = 0$  then
18:         nodeList.add( $F_j$ ,  $[\alpha^-, \alpha^+ - d]$ )
19: return nodeList

```

only moves forward while grasping, so reaching the exit range $[a_{jk}^-, a_{jk}^+]$ from the entry range turns the object by at least d , which lowers α^+ . Without holding, the robot may only leave if $d = 0$. Children with an empty interval are discarded, and children in the same T_k keep the entry range of their parent.

4) *nodeAlreadyExpanded()*: A node is skipped if an expanded node with the same key has an interval that contains its interval, at no higher cost. The key is the set for free nodes, and the set, phase and entry range for transition nodes. The check is applied when a node is taken from the queue (line 4) and before a child is added to it (line 10). Because the key and the interval describe everything that can still be reached from a node, this check is exact: a skipped node can reach nothing that the expanded one cannot reach at the same or lower cost. In CASSR, two nodes are instead treated as the same if their foothold sets lie within an empirical 2 cm of each other [14]. Unlike CASSR, we also allow the search to return to a set it has already left, since the same set can be revisited with a different interval, and dominance removes the revisits that cannot help.

5) *hasReachedTheGoal()*: A node is a goal if it lies in a free set F_i with $x_{\text{goal}} \in F_i$ and $\alpha_{\text{goal}} \in [\alpha^-, \alpha^+]$.

6) *Cost computation*: `getCostOfEdge()` returns 1 for a stroke, a return or a release, and $\mu \ll 1$ ($\mu = 10^{-3}$) for every other move. Every stroke is followed by exactly one return or release, so the cost of a plan is twice its number of grasps, plus μ times its number of set changes. As long as a plan makes fewer than $1/\mu$ set changes, minimising this cost minimises the number of grasps first and the set changes second. Releasing must cost 1, or leaving and re-entering a transition set would be a cheaper substitute for a return. [CHECK: the code also allows a return from the *entered* phase at cost 1, before any

grasp; such a move should cost μ for the cost to be exactly twice the number of grasps.] `estimateCostToGoal()` returns, with $L_{\text{max}} = \max_k L_k$,

$$h(\nu) = 2 \left\lceil \frac{\alpha^- - \alpha_{\text{goal}}}{L_{\text{max}}} \right\rceil + \mathbf{1}[\phi = \text{holding}]. \quad (3)$$

It is admissible: each remaining stroke turns the object by at most L_{max} and is followed by a return or the final release, and a node that is holding the object must still release it. It is also consistent, since no move lowers h by more than its cost. CASSR instead scales its distance-based heuristic, which makes it inadmissible and turns the search into a weighted A* [14]. With an admissible heuristic and exact dominance, the first goal node expanded has the minimum cost.

E. From sets to motions

The search fixes the sets and the grasp actions, e.g., *enter*, *stroke*, *return*, *stroke*, *release*. As in CASSR, a second stage computes the configurations, here without a solver. A backward pass from α_{goal} decides how much each visit to a transition set turns the object; it ends exactly at α_{init} because the search proved α_{goal} reachable. Waypoints are placed at door points for transit, in $F_i \cap T_k$ for entering and leaving, and on $[x_k^-, x_k^+]$ for strokes and returns. Each motion is a straight line between two waypoints of one set, so by convexity it lies in that set; we check this and the coupling for every motion. Path length is not optimised: a convex program over the chosen sets, e.g., GCS [5], could shorten the path. [TODO: implement, or keep as future work.]

F. Properties

Proposition 1 (informal): Relative to the free sets $\{F_j\}$ and the transition sets $\{T_k\}$: (i) every motion returned is a straight line inside one set that satisfies the coupling of Assumption 1, so the path is collision-free whenever the sets are; (ii) the search terminates. If a plan exists whose transit motions lie in the free sets and whose strokes lie in the transition sets, it returns one with the minimum number of grasps, with ties broken by the number of set changes. Otherwise, it returns NONE; (iii) the second stage succeeds on every sequence the search returns.

Sketch. (i) follows from convexity of the recorded sets and from the coupling check applied to every motion. (ii) Every interval end the search produces is built from the finitely many wrist values the LPs of Section IV-C return, clipped at α_{goal} , so only finitely many nodes exist per key and the search terminates. With a consistent heuristic and exact dominance, A* expands a goal node of minimum cost first [31]. Entering at the highest wrist value and leaving at the lowest wrist value of each switch edge minimises forced turning, so restricting entries and exits to these points loses no plan. (iii) The backward pass mirrors the interval updates of Algorithm 2. [CHECK: write the full proof; (iii) is asserted in our implementation but not yet proven. Overlapping transition sets would need chain-aware placement, so each set is its own segment for now.]

□

These guarantees are relative: a plan reported as optimal uses the fewest grasps *that the sets allow*. This is why the transition sets must be validated (Section VI-D).

G. Relation to a mixed-integer formulation

The same problem can be posed as a mixed-integer program, as in our earlier planner [13]: for a fixed number of grasps ℓ , choose keyframes $q_{\text{init}}, x_1, \dots, x_{2\ell}, q_{\text{goal}}$, let a binary z_{ij} indicate that keyframe i lies in set j , and increase ℓ until the program is feasible. Membership is an indicator constraint, usually in big-M form, $A_j x_i \leq b_j + M(1 - z_{ij})$. Two properties of that formulation are at odds with the sets of Section IV-C. Solvers certify a binary as integral only up to a tolerance τ , so a variable reported as $z_{ij} = 1$ may be $1 - \tau$ and relax every row of its set by $M\tau$; membership is meaningful only for sets much thicker than that, whereas transition sets are deliberately thin. And the continuous relaxation weakens as M grows, so the search tree grows with it. Tight per-row constants and extended formulations that model the union by its convex hull both help [32], at the price of a larger model, but neither removes the outer loop over ℓ . The reachability A^* has no relaxation and no integrality tolerance, and the number of grasps is an outcome of the search rather than an input.

V. EXPERIMENTAL EVALUATION

Our experiments ask four questions. Does the search return plans with the minimum number of grasps for new goals, start configurations and obstacle layouts, from a single offline stage? How does it compare with our earlier MIQP and with a sampling-based manipulation planner as the number of required grasps grows? Why must the transition sets be validated? And does validation change the plans?

A. Tasks and implementation

Both tasks turn a cap about a fixed vertical axis with a Franka Panda [33], whose wrist is restricted to $[-1, 1]$ rad, so any turn larger than 2 rad needs several grasps.

- **T1 (3-DOF)**. All arm joints but the wrist are locked: $q = (w, x_{\text{finger}}, \alpha)$, with $x_{\text{finger}} \in [-0.045, 0]$ m and $\alpha \in [-\pi, \pi]$. The decomposition has 15 regions, and the task turns the cap by 2π . **[CHECK: T1 from the MInf report (WSG gripper) or the locked Franka?]**
- **T2 (10-DOF)**. Seven arm joints, two fingers and the cap. The cap is 0.5 m in front of the robot base, with two cylindrical obstacles (radius 5 mm) next to it. The free sets come from an IRIS-NP2 clique cover with 42 regions **[CHECK: coverage target 0.98; (ϵ, δ) ; generation time]**. The reference query Q0 turns the cap from 3 to -1 rad and returns the arm to its start configuration.

The grasp constraints (1) ask the hand to be centred over the cap, within a height band, with its approach axis $z_E = {}^C R^E(x) e_z$ pointing down:

$$h(x) = [p_x, p_y, z_{E,x}, z_{E,y}]^\top, \quad (4)$$

$$g(x) = [p_z - z_{\max}, z_{\min} - p_z, z_{E,z} + c]^\top, \quad (5)$$

with $p = {}^C p^E(x)$, $[z_{\min}, z_{\max}] = [0.105, 0.110]$ m and $c = 0.9$, and the fingers held within a 1 mm band around the width of the cap. Here J_h has rank 4, leaving a six-dimensional null space. The grasp set of (2) uses $\epsilon = 10^{-3}$, $\rho = 0.05$ rad and $L_{\min} = 0.05$ rad. The wrist of the Panda satisfies the invariance of Section IV-B exactly: sweeping it over its whole range at a grasp leaves h and g unchanged to machine precision. The planner is implemented in Python with Drake [30] (v1.48). The LPs use Drake's `MathematicalProgram`, and the MIQP uses MOSEK [34]. **[TODO: hardware; code release.]**

B. Baselines and protocol

We compare against our earlier MIQP [13] (Section IV-G), with GCS [5] connecting its keyframes, and against a sampling-based planner on the manipulation graph [6], with probabilistic roadmaps for transit and transfer. **[TODO: implement the sampling-based baseline (own implementation or HPP [8]); time limit.]** For all methods, every returned motion is checked densely with Drake's collision checker. **[TODO: Run E1–E5; this is the evidence for the main claim.]** **E1, goal sweep:** α_{goal} on a grid over the cap range, with the arm configuration fixed. We report the number of grasps against the lower bound $\lceil D/L_{\max} \rceil$ for a turn D , and the search time. **E2, random queries:** N start and goal arm configurations sampled in the free sets, with random turning angles. **E3, obstacle layouts:** M placements of the obstacles, each with its own offline stage. We report offline time (decomposition, grasp set, transition sets, graph) and online time per query. **E4, baselines:** all three planners as a function of the number of grasps the query requires, on T1 and T2. We report success within the time limit, planning time and number of grasps. **E5, validation:** Q0 with and without the validation of Section IV-C.

VI. RESULTS

A. Reachability A^* on the 10-DOF task

Table I and Fig. 1 report Q0, before validation. Of the 42 free sets, 8 intersect the grasp set. Four of these transition sets span the whole wrist range ($L_k = 2$ rad), and the others span between 0.06 and 0.36 rad. The search expands 63 nodes in 4 ms and returns two strokes, one return and one release: it enters T_6 from F_2 , performs stroke, return and stroke, and releases into F_0 . Its cost, $4 + \mu$, equals the initial heuristic plus one set change, so the heuristic is tight on this query. With $D = 4$ rad and $L_{\max} = 2$ rad, two grasps is the minimum that the sets allow. The second stage turns 2 rad per stroke and produces 7 waypoints, and all 6 straight-line motions lie in their recorded set and satisfy the coupling. The online stage performs no geometric computation of its own: every intersection and extreme point it relies on was computed offline, by linear programs.

B. The trust region on the grasp set

Without the trust region ρ , the extreme points that the graph construction returns lie far outside the region where the

TABLE I
REACHABILITY A* ON T2 (10-DOF), QUERY Q0, BEFORE VALIDATION

Stage	Quantity	Value
Offline	free sets J (non-empty after slicing)	42
	transition sets K (with $L_k = 2$ rad)	8 (4)
	transit / switch edges	850 / 58
	graph construction time	[TBD]
	decomposition time	[TBD]
Online	nodes expanded	63
	search time	4 ms
	strokes / returns / releases	2 / 1 / 1
	set changes	1
	motions contained in their set	6 / 6

TABLE II
CHOOSING GRASPS WITH THE MIQP AND WITH THE REACHABILITY A*

Task	Method	Grasps	Time	Outcome
T1 (3-DOF)	MIQP	11	[TBD]	solved; GCS connected
	Reach. A*	[TBD]	[TBD]	[TBD]
T2 (10-DOF)	MIQP	–	[TBD]	numerical failure [CHECK:]
	Reach. A*	2	4 ms	solved (rel. sets)

linearisation holds: on T2 they lay 2.27 rad from the seed in one arm joint, with $\|h\|_\infty = 0.95$ and the hand 0.80 m away from the cap. With $\rho = 0.05$ rad, $\|h\|_\infty$ stays near 2×10^{-3} along the plans we report **[CHECK: numbers from the q0 run]**.

C. Comparison with the MIQP

Table II compares the two ways of choosing grasps. On T1, the MIQP returns 11 grasps after proving 10 infeasible. Each stroke turns the cap by 0.604 rad, the wrist range of the grasp set in that decomposition, and $\lceil 2\pi/0.604 \rceil = 11$. GCS then connects all 23 transit and transfer segments. On T2 it does not return a reliable solution. The numbers make the argument of Section IV-G concrete: the grasp set is $2\epsilon = 2 \times 10^{-3}$ thick, while $M = 10^6$ and MOSEK’s default integrality tolerance $\tau = 10^{-5}$ [34] give $M\tau = 10$ rad of slack on a set reported as occupied, more than any joint range. The reachability A* returns the plan above in milliseconds. **[CHECK: MIQP outcome and solve time on T2; run the reachability A* on T1 and check that it also returns 11.]**

D. Why transition sets must be validated

A planner whose guarantees are relative to its sets inherits their errors. To measure them where they matter, we swept the wrist over $[-1, 1]$ at the seed grasp, at 99 points, and compared membership in the free sets with Drake’s collision checker. The checker finds the wrist free only on $[-1, -0.19]$ and $[0.70, 0.92]$ rad **[CHECK: which bodies collide in between: fingers vs. obstacles?]**. All four decompositions we computed for T2 (IRIS-NP2 with 42 and 46 regions, and two IRIS-ZO covers with 64 regions **[CHECK: the two IRIS-ZO files may be the same cover]**) contain every point of the sweep, so all four label as collision-free the 46 configurations that

the checker rejects, and miss none of the 53 it accepts. The two-grasp plan of Fig. 1 is therefore valid relative to the sets, but collides in reality. With the longest truly free wrist interval of 0.81 rad, the minimum is $\lceil 4/0.81 \rceil = 5$ grasps: a different discrete structure, not just a different geometry.

This is not a bug in the region generators: their guarantee bounds the fraction of a region’s volume in collision, and transition sets have almost no volume (Section IV-C). The same sensitivity appeared on low-dimensional versions of the task [13], where a decomposition at 0.90 coverage contained a grasp–release sequence that GCS could not connect. For regrasp planning, what matters is not coverage of free space but the fidelity of the transition sets. **[TODO: E5: Q0 with validated transition sets; expected 5 grasps with no change to the search.]**

E. Generalisation across goals, starts and scenes

[TODO: Results of E1–E4 (Section V-B).]

VII. DISCUSSION AND LIMITATIONS

a) *Validating where it matters:* The weakest link is not the search but the sets it searches over (Section VI-D). A dense check is not a certificate; certifying the thin transition sets with C-IRIS [2], or seeding the decomposition on the grasp set, would strengthen the guarantees without changing the search.

b) *Constructing grasp sets:* We take two ideas from NLP sampling [15], its Gauss–Newton phase and its null-space view of the constraints, but not its interior samplers, which we first tried and could not turn into usable polytopes. Seeding IRIS regions with grasp samples produced many tiny polytopes, because a region grown around a thin, curved set is bounded by its curvature in every direction. Fitting a union of polytopes to the samples with a small network, each neuron a half-space, a region the maximum over its half-spaces and the set the minimum over regions, never fitted the grasp set closely enough to plan with. Hit-and-run itself mixed poorly in the wrist and cap coordinates, since its line search shares one step size across all of them. What works instead is to exploit the structure of the constraints, above all their exact invariance to the wrist, rather than to rediscover it from samples. The price is locality: one seed yields one grasp set, valid near that seed. The restarts of NLP sampling spread seeds across the components of the grasp manifold, and are the natural way to obtain the several grasp sets the search already accepts. **[TODO: numbers for the attempts, if space allows.]**

c) *Scope:* Exact intervals rely on Assumption 1, which covers caps, knobs, valves and cranks but not general pick-and-place. With more object coordinates, the object configurations reachable in a set become sets that could be propagated with Minkowski sums and intersections, as in CASSR [14], at the cost of exact dominance and with nonlinear grasp transforms to approximate.

d) *Optimality and offline cost:* The search minimises the number of grasps, then the number of set changes, which is only a weak proxy for distance travelled. As in CASSR, the sets are fixed before the configurations, so path length is not

optimal. A new obstacle layout requires a new decomposition, which dominates the offline cost: in a separate run, IRIS-NP2 took 120 s for 25 regions at a 0.98 coverage target **[CHECK: different settings]**. GPU implementations [29] reduce this cost, and new start and goal queries reuse the graph unchanged.

VIII. CONCLUSION

We represented the manipulation graph of multi-regrasp tasks with convex sets, so that the switches between moving alone and moving the object are convex transition sets, and searched them with an A* whose nodes pair a set with the interval of object configurations reachable inside it. For objects with one coordinate coupled to the wrist, the search is complete and returns the fewest grasps the sets allow, without fixing that number in advance and without integer variables. The search thus carries over to regrasping an idea matured in footstep planning, that of searching over convex reachable sets rather than individual contacts, and extends it to contact sets that have to be constructed and to contacts that move the object. Its guarantees are only as good as the transition sets are faithful near contact, which is why we validate them. **[TODO: one sentence on E1–E5 once run.]**

ACKNOWLEDGMENT

[TODO: acknowledgments; funding is in the title footnote.]

REFERENCES

- [1] R. Deits and R. Tedrake, “Computing large convex regions of obstacle-free space through semidefinite programming,” in *Algorithmic Foundations of Robotics XI*. Springer, 2015, pp. 109–124.
- [2] H. Dai, A. Amice, P. Werner, A. Zhang, and R. Tedrake, “Certified polyhedral decompositions of collision-free configuration space,” *The International Journal of Robotics Research*, vol. 43, no. 9, pp. 1322–1341, 2024.
- [3] M. Petersen and R. Tedrake, “Growing convex collision-free regions in configuration space using nonlinear programming,” *arXiv preprint arXiv:2303.14737*, 2023.
- [4] P. Werner, A. Amice, T. Marcucci, D. Rus, and R. Tedrake, “Approximating robot configuration spaces with few convex sets using clique covers of visibility graphs,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2024.
- [5] T. Marcucci, M. Petersen, D. von Wrangel, and R. Tedrake, “Motion planning around obstacles with convex optimization,” *Science Robotics*, vol. 8, no. 84, p. eadf7843, 2023.
- [6] T. Siméon, J.-P. Laumond, J. Cortés, and A. Sahbani, “Manipulation planning with probabilistic roadmaps,” *The International Journal of Robotics Research*, vol. 23, no. 7–8, pp. 729–746, 2004.
- [7] K. Hauser and J.-C. Latombe, “Multi-modal motion planning in non-expansive spaces,” *The International Journal of Robotics Research*, vol. 29, no. 7, pp. 897–915, 2010.
- [8] J. Mirabel, S. Tonneau, P. Fernbach *et al.*, “HPP: A new software for constrained motion planning,” in *Proc. IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2016, pp. 383–389.
- [9] W. Wan and K. Harada, “Regrasp planning using 10,000s of grasps,” in *Proc. IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2017, pp. 1929–1936.
- [10] P. Lertkultanon and Q.-C. Pham, “A certified-complete bimanual manipulation planner,” *IEEE Transactions on Automation Science and Engineering*, vol. 15, no. 3, pp. 1355–1368, 2018.
- [11] M. Toussaint, “Logic-geometric programming: An optimization-based approach to combined task and motion planning,” in *Proc. International Joint Conference on Artificial Intelligence (IJCAI)*, 2015, pp. 1930–1936.
- [12] S. Levit and M. Toussaint, “Regrasp maps for sequential manipulation planning,” *arXiv preprint arXiv:2507.12407*, 2025.
- [13] J. López Gómez, “Optimisation-based manipulation planning in convex decompositions of C-free,” MInf Project Report, School of Informatics, University of Edinburgh, 2025.
- [14] J. Wang and S. Tonneau, “CASSR: Continuous A-star search through reachability for real time footstep planning,” *arXiv preprint arXiv:2603.02989*, 2026.
- [15] M. Toussaint, C. V. Braun, and J. Ortiz-Haro, “NLP sampling: Combining MCMC and NLP methods for diverse constrained sampling,” *arXiv preprint arXiv:2407.03035*, 2024.
- [16] P. Tournassoud, T. Lozano-Pérez, and E. Mazer, “Regrasping,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 1987, pp. 1924–1928.
- [17] R. Alami, T. Siméon, and J.-P. Laumond, “A geometrical approach to planning manipulation tasks. the case of discrete placements and grasps,” in *The Fifth International Symposium on Robotics Research*. MIT Press, 1990, pp. 453–463.
- [18] P. S. Schmitt, W. Neubauer, W. Feiten, K. M. Wurm, G. v. Wichert, and W. Burgard, “Optimal, sampling-based manipulation planning,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2017.
- [19] Z. Kingston, M. Moll, and L. E. Kavraki, “Sampling-based methods for motion planning with constraints,” *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 1, pp. 159–185, 2018.
- [20] C. R. Garrett, R. Chitnis, R. Holladay, B. Kim, T. Silver, L. P. Kaelbling, and T. Lozano-Pérez, “Integrated task and motion planning,” *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 4, pp. 265–293, 2021.
- [21] S. Liu, T. Zhao, B. P. Graesdal *et al.*, “Approximately optimal global planning for contact-rich SE(2) manipulation on a graph of reachable sets,” *arXiv preprint arXiv:2601.10827*, 2026.
- [22] R. J. Griffin, G. Wiedebach, S. McCrory, S. Bertrand, I. Lee, and J. Pratt, “Footstep planning for autonomous walking over rough terrain,” in *Proc. IEEE-RAS International Conference on Humanoid Robots (Humanoids)*, 2019, pp. 9–16.
- [23] R. Deits and R. Tedrake, “Footstep planning on uneven terrain with mixed-integer convex optimization,” in *Proc. IEEE-RAS International Conference on Humanoid Robots (Humanoids)*, 2014, pp. 279–286.
- [24] S. Tonneau, D. Song, P. Fernbach, N. Mansard, M. Taïx, and A. Del Prete, “SL1M: Sparse L1-norm minimization for contact planning on uneven terrain,” in *Proc. IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 6604–6610.
- [25] T. Marcucci, J. Umenberger, P. Parrilo, and R. Tedrake, “Shortest paths in graphs of convex sets,” *SIAM Journal on Optimization*, vol. 34, no. 1, pp. 507–532, 2024.
- [26] B. P. Graesdal, S. Y. C. Chia, T. Marcucci *et al.*, “Towards tight convex relaxations for contact-rich manipulation,” in *Proc. Robotics: Science and Systems (RSS)*, 2024.
- [27] S. Y. C. Chia, R. H. Jiang, B. P. Graesdal, L. P. Kaelbling, and R. Tedrake, “GCS*: Forward heuristic search on implicit graphs of convex sets,” in *Workshop on the Algorithmic Foundations of Robotics (WAFR)*, 2024.
- [28] P. Werner, T. Cohn, R. H. Jiang *et al.*, “Faster algorithms for growing collision-free convex polytopes in robot configuration space,” in *International Symposium on Robotics Research (ISRR)*, 2024.
- [29] P. Werner, R. Cheng, T. Stewart, R. Tedrake, and D. Rus, “Superfast configuration-space convex set computation on GPUs for online motion planning,” *arXiv preprint arXiv:2504.10783*, 2025.
- [30] R. Tedrake and the Drake Development Team, “Drake: Model-based design and verification for robotics,” <https://drake.mit.edu>, 2019.
- [31] P. E. Hart, N. J. Nilsson, and B. Raphael, “A formal basis for the heuristic determination of minimum cost paths,” *IEEE Transactions on Systems Science and Cybernetics*, vol. 4, no. 2, pp. 100–107, 1968.
- [32] J. P. Vielma, “Mixed integer linear programming formulation techniques,” *SIAM Review*, vol. 57, no. 1, pp. 3–57, 2015.
- [33] S. Haddadin, S. Parusel, L. Johannsmeier *et al.*, “The Franka Emika robot: A reference platform for robotics research and education,” *IEEE Robotics & Automation Magazine*, vol. 29, no. 2, pp. 46–64, 2022.
- [34] MOSEK ApS, “The MOSEK optimization toolbox,” <https://www.mosek.com>, 2024.